



# AuditBench KTO (Llama-3.3-70B; 2 variants)


**$\mu$ -decisiveness · fitted preference strength**




**P(repeat) · self-agreement**




**P(agree | A,B vs B,A) · order**



**P(no cycle) · transitivity**



**P(agree | Q<sub>i</sub> vs Q<sub>j</sub>) · framing**



Llama 3.x baselines (size →)
  base for fine-tuning
  AR = Anti-AI regulation
  AW = Animal welfare
  base model performance
  chance floor