



AuditBench KTO (Llama-3.3-70B; 3 variants)


μ -decisiveness · fitted preference strength




P(repeat) · self-agreement




P(agree | A,B vs B,A) · order



P(no cycle) · transitivity



P(agree | Q_i vs Q_j) · framing



Llama 3.x baselines (size →)
 base for fine-tuning
 AR = Anti-AI regulation
 AW = Animal welfare
 base model performance
 chance floor