

Gemma-3-27B-it: coherence vs quantization precision
(bf16 → int8 → nf4; dashed = bf16 level; y-axes zoomed — note differences are tiny)

