

Consistency (agreement-probability) vs model size
 Gemma-3 · Qwen2.5 · Llama-3.x + GPT-OSS 20B/120B (markers = budget)

