

OLMo-2-7B: coherence vs pretraining compute
(flat through pretraining; jumps only after instruction tuning ★)

