

Sentiment-coherence metrics vs model size (all benchmarks)
 Gemma-3 · Qwen2.5 · Llama-3.x + GPT-OSS-20B (markers = reasoning budget, at 20B)

